

# Project Report: Multilingual Transla- tor

*Architecture, Setup, and Usage Guide*

Technical Documentation



# Contents

<b>1</b>	<b>Project Overview &amp; Architecture</b>	<b>1</b>
1.1	Introduction . . . . .	1
1.2	Problem & Solution . . . . .	1
1.3	System Design . . . . .	1
1.4	Model Decisions . . . . .	2
<b>2</b>	<b>Installation &amp; Setup</b>	<b>3</b>
2.1	Prerequisites . . . . .	3
2.2	Step-by-Step Installation . . . . .	3
2.2.1	1. Clone the Repository . . . . .	3
2.2.2	2. Create a Virtual Environment . . . . .	4
2.2.3	3. Install Dependencies . . . . .	4
2.3	GPU Configuration (Optional but Recommended) . . . . .	4
2.4	Directory Structure Verification . . . . .	5
2.5	Troubleshooting Common Issues . . . . .	5
<b>3</b>	<b>Usage Manual</b>	<b>6</b>
3.1	Running the Application . . . . .	6
3.1.1	Startup Logs . . . . .	6
3.2	Interface Overview . . . . .	6
3.3	Step-by-Step Workflow . . . . .	7
3.4	Example Scenarios . . . . .	7
3.4.1	Scenario 1: News Article (Hindi) . . . . .	7
3.4.2	Scenario 2: Tech Review (French) . . . . .	8
3.5	Advanced Customization . . . . .	8
3.5.1	Using a Smaller Model . . . . .	8
3.5.2	Changing the Summary Style . . . . .	8
<b>4</b>	<b>Code Implementation Details</b>	<b>9</b>
4.1	Translation Service . . . . .	9
4.1.1	Handling Multilingual Inputs . . . . .	9
4.2	Summarization Service . . . . .	10
4.3	The Application Controller . . . . .	11

# Chapter 1

# Project Overview & Architecture

## 1.1 Introduction

In today's interconnected world, language barriers remain a significant challenge. The **Multilingual Translator & Summariser** is an AI-powered application designed to bridge this gap. It allows users to input text in *any* language, automatically translates it into English, and then provides a concise summary of the content.

## 1.2 Problem & Solution

- **The Problem:** Information Overload. Valuable documents (news, research, user feedback) are often available only in local languages and are too long to digest quickly.
- **The Solution:** A single-click pipeline.
  1. **Translate:** Convert local language 'X' to English using NLLB-200.
  2. **Summarize:** Condense the English text using BART-Large.
  3. **Present:** Display both outputs in a clean, reactive Web UI.

## 1.3 System Design

The system follows a modular architecture, separating the core AI logic from the user interface.

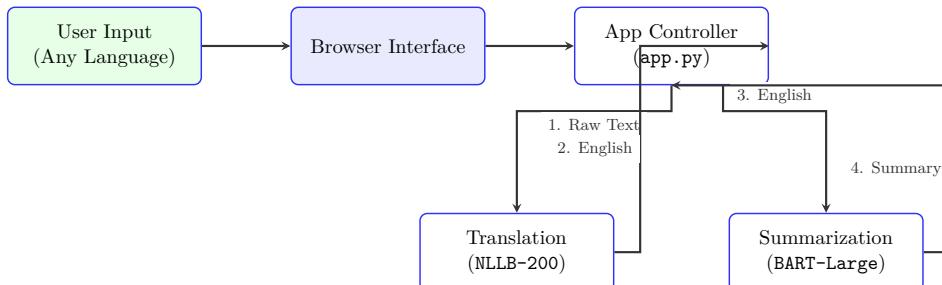


Figure 1.1: Data Flow Architecture

## 1.4 Model Decisions

**Translation Model:** We selected ‘facebook/nllb-200-distilled-600M’.

- **Pros:** Supports 200+ languages, high accuracy, reasonable size (1.2GB).
- **Why Distilled?** The full model is 54GB, which is impossible to run on consumer hardware. The distilled version offers 90% of the performance at 2% of the size.

**Summarization Model:** We selected ‘facebook/bart-large-cnn’.

- **Pros:** State-of-the-art for abstractive summarization on news articles.
- **Why CNN?** It was fine-tuned on the CNN/DailyMail dataset, making it excellent for understanding journalistic and factual content.

# Chapter 2

# Installation & Setup

This guide assumes you are running on Windows, but the commands are similar for Mac/Linux. We will cover the installation of Python, the creation of a virtual environment, and the installation of the machine learning dependencies.

## 2.1 Prerequisites

Before diving in, ensure you have the following installed on your system:

1. **Python 3.8+:** Hugging Face libraries require a modern Python version.  
Verification: `python --version`.
2. **Git:** Standard for version control, useful for cloning the repository.
3. **Internet Connection:** To download model weights. The models (NLLB-200 and BART-Large) are approx 2-3 GB combined.
4. **Hardware:** A GPU (NVIDIA GTX 1650 or better) is recommended for fast inference. If you use a CPU, expect delays of 10-15 seconds per request.

## 2.2 Step-by-Step Installation

### 2.2.1 1. Clone the Repository

Open your terminal (PowerShell or Command Prompt) and navigate to your workspace.

```
1 # Clone the project code
2 git clone https://github.com/your-repo/multilingual-
   translator.git
3
4 # Enter the directory
```

```
5 cd multilingual-translator
```

### 2.2.2 2. Create a Virtual Environment

It is "best practice" to isolate dependencies so they don't conflict with other projects.

```
1 # Create the environment named 'venv'
2 python -m venv venv
3
4 # Activate it (Windows)
5 # Your prompt should change to start with (venv)
6 venv\Scripts\activate
7
8 # Activate it (Linux/Mac)
9 source venv/bin/activate
```

### 2.2.3 3. Install Dependencies

We rely on `torch`, `transformers`, and `gradio`.

```
1 # Upgrade pip first to avoid errors
2 pip install --upgrade pip
3
4 # Install dependencies from the file
5 pip install -r requirements.txt
```

## 2.3 GPU Configuration (Optional but Recommended)

Standard `pip install torch` usually installs the CPU version. To enable GPU support:

1. Check your CUDA version (open NVIDIA Control Panel or run `nvidia-smi`).
2. Go to <https://pytorch.org/get-started/locally/>.
3. Copy the command for your version. For example (CUDA 11.8):

```
1 pip install torch torchvision torchaudio --index-url
  https://download.pytorch.org/whl/cu118
```

## 2.4 Directory Structure Verification

After installation, your project folder should look like this:

```
1 multilingual-translator/
2 |-- venv/                      # Virtual Environment (do not
3 |   touch)
4 |-- src/
5 |   |-- __init__.py
6 |   |-- translation.py    # Translation Logic
7 |   |-- summarization.py # Summarization Logic
8 |-- app.py                     # Main Entry Point
9 |-- requirements.txt # Dependency List
9 |-- README.md                 # Quick Start Guide
```

### Warning

Do not commit the `venv/` folder to Git! It is specific to your machine and very large. Use a `.gitignore` file to exclude it.

## 2.5 Troubleshooting Common Issues

- **Error: "Module not found"**: Ensure you activated the virtual environment before running the app.
- **Error: "Torch not compiled with CUDA enabled"**: You installed the CPU version of Torch. Uninstall it (`pip uninstall torch`) and reinstall the GPU version using the command above.
- **Slow Performance**: If the translation takes >30 seconds, you are likely running on CPU. This is expected for large models like NLLB-200.
- **Download Fails**: These models are hosted on Hugging Face Hub. If you have a firewall, you might need to use a proxy or check your internet connection.

# Chapter 3

# Usage Manual

This chapter guides you through running the application and understanding the user interface.

## 3.1 Running the Application

With your environment active, run the following command in your terminal:

```
1 python app.py
```

### 3.1.1 Startup Logs

You should see output similar to the following. Note the model loading times:

```
Initializing AI Services...
Loading Translation Model: facebook/nllb-200-distilled-600M...
(This may take 10-20 seconds on first run to download 1.2GB)
Loading Summarization Model: facebook/bart-large-cnn...
(This may take 10-15 seconds to download 1.6GB)
Services loaded successfully.
Running on local URL: http://127.0.0.1:7860
```

Open your browser (Chrome/Edge/Firefox) and navigate to <http://127.0.0.1:7860>.

## 3.2 Interface Overview

The User Interface (UI) is built with Gradio and consists of three main sections:

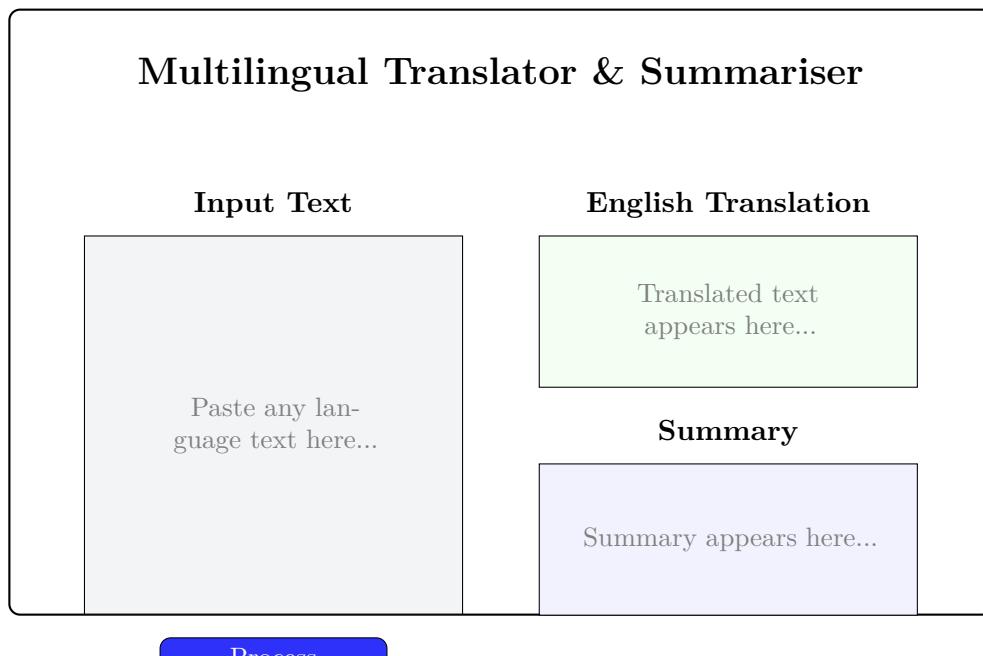


Figure 3.1: UI Layout Mockup

### 3.3 Step-by-Step Workflow

1. **Input Field:** Copy and paste any text into the large box on the left. The model handles over 200 languages (Spanish, Hindi, French, Japanese, etc.).
2. **Process Button:** Click the "Process" button. The button will show a loading spinner while the GPU processes the request.
3. **Outputs:**
  - **Top Right (Translation):** The system first translates the input into English. This preserves the full detail of the original text.
  - **Bottom Right (Summary):** The system then reads the English translation and creates a 2-3 sentence abstractive summary.

### 3.4 Example Scenarios

#### 3.4.1 Scenario 1: News Article (Hindi)

- **Input:** "Bharat ek vishaal desh hai..." (A long paragraph about India's geography).

- **Action:** User clicks Process.
- **Translation:** "India is a vast country..."
- **Summary:** "India is geographically diverse with ancient culture."

### 3.4.2 Scenario 2: Tech Review (French)

- **Input:** "Ce nouvel ordinateur portable est incroyablement rapide mais la batterie..."
- **Translation:** "This new laptop is incredibly fast but the battery..."
- **Summary:** "The laptop has excellent performance but poor battery life."

#### Warning

The first request might be slow (10-20 seconds) as the computer "warms up" the models (JIT Compilation). Subsequent requests will be much faster.

## 3.5 Advanced Customization

Want to change the models? You can modify 'src/translation.py':

### 3.5.1 Using a Smaller Model

If you are on an older laptop, switch to the 600M distilled version or even the 1.3B version if you have more RAM.

```
1 # In src/translation.py
2 self.model_name = "facebook/nllb-200-distilled-600M" #
3 # self.model_name = "facebook/nllb-200-3.3B"           #
# Very Slow, More Accurate
```

### 3.5.2 Changing the Summary Style

You can make the summary bullet-pointed by editing 'src/summarization.py':

```
1 # In src/summarization.py
2 def summarize(self, text):
3     prompt = f"Summarize this in bullet points: {text}"
4     # ... rest of code
```

# Chapter 4

## Code Implementation Details

This chapter provides a deep dive into the implementation. We use object-oriented programming to keep the codebase modular and testable.

### 4.1 Translation Service

File: `src/translation.py`

The ‘TranslationService‘ class encapsulates the complexity of the NLLB model.

```
1  class TranslationService:
2      def __init__(self, model_name="facebook/nllb-200-
3          distilled-600M"):
4          # We check for GPU availability automatically
5          self.device = 0 if torch.cuda.is_available() else
6              -1
7
8          print(f"Loading Translation Model: {model_name
9              }...")
10         self.tokenizer = AutoTokenizer.from_pretrained(
11             model_name)
12         self.model = AutoModelForSeq2SeqLM.
13             from_pretrained(model_name)
14
15         if self.device == 0:
16             self.model = self.model.to("cuda")
```

Listing 4.1: Initializing NLLB

#### 4.1.1 Handling Multilingual Inputs

The core challenge with NLLB is that it requires a ”Forced BOS Token“ (Beginning of Sentence) to know which language to translate *into*.

```

1  def translate(self, text, target_lang="eng_Latn"):
2      if not text: return ""
3
4      # 1. Tokenize Input
5      inputs = self.tokenizer(text, return_tensors="pt"
6                               , padding=True)
7      if self.device == 0:
8          inputs = {k: v.to("cuda") for k, v in inputs.
9                     items()}
10
11     # 2. Force Target Language Token
12     # "eng_Latn" is the NLLB code for English
13     forced_bos = self.tokenizer.convert_tokens_to_ids
14         (target_lang)
15
16     # 3. Generate
17     with torch.no_grad():
18         tokens = self.model.generate(
19             **inputs,
20             forced_bos_token_id=forced_bos,
21             max_length=512
22         )
23
24     # 4. Decode
25     return self.tokenizer.batch_decode(tokens,
26                                         skip_special_tokens=True)[0]

```

Listing 4.2: Forcing English Output

## 4.2 Summarization Service

File: `src/summarization.py`

The summarizer is simpler because we leverage the Hugging Face ‘pipeline’ abstraction.

```

1  class SummarizationService:
2      def __init__(self):
3          # Pipelines handle tokenization, device placement
4          # , and decoding internally
5          self.summarizer = pipeline("summarization", model
6                                     ="facebook/bart-large-cnn")
7
8      def summarize(self, text):
9          # We set tight constraints on length to ensure
10             concise outputs
11         result = self.summarizer(
12             text,
13             max_length=130,    # Max summary length

```

```
11         min_length=30,      # Min summary length
12         do_sample=False    # Deterministic output
13     )
14     return result[0]['summary_text']
```

### 4.3 The Application Controller

File: `app.py`

Finally, we tie everything together.

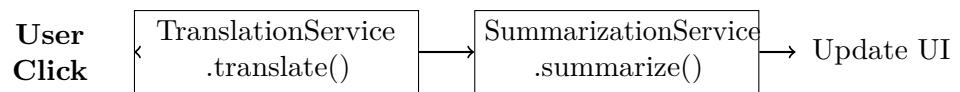


Figure 4.1: Execution Flow

The `app.py` script initializes these services once (globally) so they don't reload on every request, which would be very slow.